Certainly! Here's the comprehensive table you requested. It includes:

Model	Core Ideas (Pros)	Limitations (Cons)	Technical Requirements	Time Encoding	Forecast Target	Eval Metrics	Method
Prophet	Seasonality & holiday built-in	Weak with non-linearity	Minimal	Internal	Price, seasonality	RMSE, MAPE	Additive Model
LSTM	Captures temporal non-linearity	Needs much data	No stationarity	Cyclical or embedded	Price, return, volatility	RMSE, MAPE, R ²	DL (RNN)
GRU	Efficient vs. LSTM	Misses long dependencies	Same as LSTM	Same	Price, volatility	Same as LSTM	DL (RNN)
Bi-LSTM	Considers future/past context	Future not always known	Same as LSTM	Same	Price (post-event), sentiment	Same as LSTM	DL (RNN)
CNN (Time Series)	Local pattern learning	Poor for long term	No stationarity	Treat time as channel	Price trend, short-term	MSE, Accuracy (patterns)	DL (CNN)
Transformer	Global attention, long memory	Heavy compute	None, large data needed	Sin/Cos positional	Price, return, volatility, regime	RMSE, MAE, R ²	DL (Attention)
Informer/Autoformer	Long horizon attention models	Complex	No stationarity	Learned embeddings	Price, volatility, regime	RSE, MASE	DL (Efficient Transformer)
N-HiTS	Fast + scalable forecasting	Research stage	None	Internal	Trend, price	MASE, RSE	DL (Residual Hierarchical)
DeepAR	Probabilistic & multiple series	Needs AR-like structure	None	Internal embeddings	Price, volatility	NLL, CRPS	DL (Auto-regressive RNN)
Deep State Space	Uncertainty + deep learning	Heavy compute	None	Fourier/time basis	Regime, macro-volatility	NLL, RMSE	Bayesian DL
TFT	Interpretable, deep attention	Tuning complex	None	Learned embeddings	Price, volatility, return	MAE, RMSE, R ²	DL (Attention + Gating)
XGBoost/LightGBM	Strong on structured data	Manual lags needed	Scaling + feature prep	One-hot / cyclical	Price, return, volatility	R ² , MAE, RMSE	Tree Boosting
TabNet	Deep learning for tabular data	New; needs fine-tuning	No stationarity	Embedded internally	Price, return	RMSE, R ²	DL (Attentive FCN)
DDPG (RL)	Continuous actions	Reward shaping is tricky	Agent tuning	In state vector	Return, trading actions	Sharpe, PnL, DD	RL (Policy Gradient)
TD3	Stable vs. DDPG (2 critics)	Slower training	Same	Same	Return, risk mgmt	Sharpe, PnL	RL (Value + Policy)
SAC	High exploration entropy	Hyperparameter sensitive	High compute	Same	Risk-adjusted returns	Sharpe, Entropy, Profit	RL (Entropy Optim.)
PPO	Stable & safe RL training	Reward design matters	Sample efficiency	Same	Portfolio weights	Sharpe, MDD	RL (Clipped Policy)
A3C	Fast, parallel training	Sample inefficiency	Same	Same	Return, policy learning	Sharpe, Reward curve	RL (Actor-Critic)
AlphaStock	Combines RL + DL like AlphaZero	Very compute heavy	RL + DL infra	Structured state	Portfolio strategy	Profit, Sharpe	Hybrid (Policy RL + DL)
Meta-Learning (MAML, etc.)	Fast adaptation to change	Experimental	Time-aware adaptation	Dynamic embeddings	Regime change, multi-goal	Transfer performance	Meta-Learning



Project-Wide Notes and Critical Considerations

1. Stationarity Concerns

- Classical models (ARIMA, SARIMA) require stationarity (use differencing, detrending).
- DL/RL models (LSTM, Transformer, PPO, SAC, etc.) do not require stationarity, but large shifts (structural breaks) can still harm them.

Best practice: Detrend input data (optional) and scale features.

2. Feature Encoding for Time

• Cyclical Encoding (preferred for DL/RL): month_sin=sin(2π×month12),month_cos=cos(2π×month12)month_sin=sin(2π×12month),month_cos=cos(2π×12month)

• One-hot encoding or embedding layers for categorical time features (week, day-of-week) when used in tree models or DL hybrid models.

3. Technical Indicators to Include (examples)

Type Examples

Trend Moving Averages (SMA, EMA), MACD

Momentum RSI, Stochastic Oscillator, Williams %R

Volatility ATR (Average True Range), Bollinger Bands

Volume-based OBV (On-Balance Volume), Volume Price Trend

4. Fundamental Indicators

- Earnings per Share (EPS)
- Price-to-Earnings Ratio (P/E)
- Price-to-Book Ratio (P/B)
- Dividend Yield
- Debt-to-Equity Ratio (D/E)
- Revenue, Net Income Growth
- Free Cash Flow
- Insider Buying Activity

5. Indices and Macro Factors to Monitor

- S&P500, NASDAQ, Dow Jones (US indices)
- VIX (Volatility Index)
- Interest Rates (e.g., Fed Funds Rate)
- Inflation Data (CPI, PPI)
- Employment Data (Unemployment rate)
- Global Economic Sentiment Indices
- Oil Prices, Gold Prices, etc.

6. Time Series Specific Concerns

Concern Description

Seasonality Holiday effects, quarterly earnings

Regime Changes Market crash, major political events

Noise Random market movements unrelated to fundamentals

Outliers Extreme points can distort learning

✓ Take Care Notes (Golden Rules)

- Always normalize/standardize inputs (especially for DL).
- Watch for data leakage (future information contaminating the past).
- **Expand features** carefully: lags, rolling means, cumulative sums.
- Use cross-validation specific to time series (e.g., expanding windows).
- Regular model retraining is needed (market changes constantly).

Common Forecasting Goals by Model

Forecast Target Suggested Models

Price ARIMA, LSTM, Transformer, DeepAR, TFT

LSTM, GRU, RL (DDPG, PPO), TFT Return

Transformer, DeepAR, SAC Volatility

CNN, LSTM, XGBoost Momentum

PPO, SAC, TD3, TFT Risk

Portfolio Allocation AlphaStock, RL methods (DDPG, PPO, A3C)

Multi-horizon / Sequence Output Transformer, TFT, N-HiTS, DeepAR, Meta-Learning

📚 Notes on Evaluation Metrics

Metric Purpose

RMSE / MAE / MSE General accuracy on price prediction

MAPE / SMAPE Scaled error, interpretability

Sharpe Ratio Return/risk in trading

Max Drawdown (MDD) Risk of capital loss

CRPS / NLL Probabilistic forecasts (DeepAR, Bayesian)

R² (R-squared) Goodness of fit

PnL (Profit) Total return for RL agents

- Data flow and logic steps for each model.
- In a way that teaches you how to implement a real stock market prediction system with each approach.
- Focused on Deep Learning (DL) and Reinforcement Learning (RL).

I'll break it down carefully.



✓ Stock Prediction - Model Implementation Flows



4 1. Deep Learning Models (DL)

■ LSTM / GRU for Stock Prediction

```
Raw Data (OHLCV + indicators + time features)
Data Cleaning (remove NaN, outliers)
Feature Engineering
    - Technical indicators (RSI, MACD)
    - Cyclical encoding (month/day/week sin-cos)
    - Normalize/standardize inputs
Data Preparation
    - Create sequences (lookback windows)
    - X = past N days' features
    - y = future price / return
Train-Test Split (time-aware split)
Build Model
    - LSTM layers (1-3 layers)
    - Dropout (for regularization)
    - Dense output (linear or sigmoid)
Train Model (early stopping, low LR)
Evaluate
    - RMSE, MAE, R<sup>2</sup> on validation set
Predict future price/return
```

Bi-LSTM (Bidirectional LSTM)

(Same as above, **BUT**)

Use Bidirectional wrapper around LSTM layers to allow backward+forward temporal attention.

CNN for Stock Prediction

```
Java
CopyEdit
Raw Data (candlestick OHLCV, indicators)

↓
Data Cleaning & Feature Engineering

↓
Data Shaping

— Convert into 2D shape (sequence length × features)

↓
Build CNN Model

— Conv1D or Conv2D layers

— Pooling layers (optional)

— Flatten → Dense output

↓
Train and Evaluate

↓
Predict patterns (short-term momentum, price)
```

Transformer (Vanilla)

```
mathematica
CopyEdit
Raw Data (OHLCV, Indicators, Macro data)

Feature Engineering

- Technical + Fundamental data
- Cyclic Encode time features

Build Input Sequences

Positional Encoding

Build Transformer Encoder-Decoder
- Multi-Head Attention
- Feedforward Networks

Train Model

Evaluate
- RMSE, MAE, MAPE

Predict multiple steps ahead
```

□ Informer / Autoformer / Reformer

(Almost same as Transformer, but optimized)

java CopyEdit

```
Raw Data

Feature Engineering

Long-Horizon Sequence Preparation

Efficient Transformer Variant

- Sparse attention

- Decomposition of trends (Autoformer)

Train and Forecast long horizons
```

DeepAR (Amazon)

```
java
CopyEdit
Raw Data (multiple stock series)

Time Series Normalization

↓
Feed Past Data (Autoregressive Inputs)

↓
Recurrent Neural Network (RNN)

↓
Output Future Probabilistic Forecasts

- Predict mean and variance
```

Temporal Fusion Transformer (TFT)

```
sql
CopyEdit
Raw Data

Static Features + Dynamic Features Separation

Embedding Layers

- Time Embeddings

- Feature Embeddings

Gated Residual Networks (GRN)

Attention Layers

- Select important historical steps

- Select important features

Train with Quantile Loss

Forecast future price/volatility/return
```

2. Reinforcement Learning Models (RL)

DDPG / TD3 / SAC / PPO (RL for Trading)

```
CopyEdit
Raw Data (OHLCV, indicators, macro)
Environment Setup
    - Define State (features, indicators, past returns)
    - Define Action (buy, sell, hold, allocation %)
    - Define Reward (portfolio return, Sharpe Ratio)
Agent Setup
    - Actor network: outputs action
    - Critic network: evaluates action
Training Loop
    - Observe State
    - Select Action
    - Execute Action
    - Receive Reward
    - Update Policy (Actor + Critic)
   \downarrow
Evaluation
    - Total return
    - Sharpe Ratio
    - Max Drawdown
```

■ AlphaStock (Advanced RL + DL approach)

```
- Actor network
- Critic network

↓

Self-play / Simulated Trading

↓

Reward Shaping
- Profit, Risk Control, Transaction Costs

↓

Training Massive Simulations
↓

Evaluate on Unseen Market Conditions
```

Meta-Learning (MAML for Stock Prediction)

```
csharp
CopyEdit
Raw Data (market tasks)

Meta-Training Phase

- Learn an initialization that adapts fast

Meta-Testing Phase

- Fine-tune quickly on new stocks / new market regimes

Adapted Model predicts price, return, regime
```



Overall Deep Learning/RL Stock Prediction Workflow

markdown CopyEdit

- 1. Problem Formulation
 - Price prediction? Return? Volatility?
- 2. Data Collection
 - OHLCV, Technical, Macro, Fundamental
- 3. Data Cleaning
 - NaN removal, Outlier smoothing
- 4. Feature Engineering
 - Technical Indicators
 - Cyclical Date Encoding
- 5. Model Selection
 - DL (LSTM, Transformer) for pure forecasting
 - RL (PPO, DDPG) for action/trading strategy
- 6. Model Training
 - Time-aware validation (expanding window CV)
- 7. Model Evaluation
 - RMSE, MAPE, Sharpe Ratio
- 8. Backtestind
 - Simulate strategies on historical data
- 9. Deployment
 - Predict live or trade live
- Monitoring

- Retrain and adapt to market regime changes



- **DL = Prediction** (future values)
- RL = Decision Making (how to act in market)

They can work together (first predict price, then act accordingly)!

Multicollinearity (high correlation between two or more input features) can seriously hurt many models (especially linear ones like ARIMA, SARIMAX, or even tree models if too severe).

Here are some excellent Python libraries and methods to detect multicollinearity:



Libraries and Methods to Check Multicollinearity

Library	Functionality	Example Function / Method	Notes
statsmodels	Calculate Variance Inflation Factor (VIF)	variance_inflation_factor()	Gold standard for VIF checks
pingouin	VIF and partial correlations	pingouin.vif()	Very easy, returns clean DataFrame
scikit-learn	Correlation matrix + PCA analysis	np.corrcoef(), PCA()	Not direct VIF but useful
seaborn + matplotlib	Visual correlation heatmaps	<pre>sns.heatmap(corr_matrix)</pre>	Fast visual spotting of collinear pairs
pycaret	Built-in multicollinearity removal in preprocess	<pre>setup(, remove_multicollinearity=True)</pre>	Full pipeline automation
feature-engine	Select/remove collinear features	<pre>feature_engine.selection.DropCorrel atedFeatures</pre>	Specialized for data cleaning



Typical Code Snippets for Checking Multicollinearity:

1. VIF Calculation (statsmodels)

from statsmodels.stats.outliers_influence import variance_inflation_factor import pandas as pd # Assume `X` is your features DataFrame vif_data = pd.DataFrame() vif_data["feature"] = X.columns

```
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif_data)
```

- **VIF > 5**: Moderate multicollinearity
- **VIF > 10**: Serious multicollinearity

2. Correlation Matrix with Heatmap (seaborn)

```
python
CopyEdit
import seaborn as sns
import matplotlib.pyplot as plt
corr = X.corr()
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.show()
```

• Look for correlation coefficients > 0.8 or < -0.8 between features!

3. Automatic VIF Check (pingouin)

CopyEdit import pingouin as pg vif = pg.vif(X) print(vif)

Much faster, good for big DataFrames.



Best Practice for DL models:

- Deep Learning (LSTM, Transformer, etc.) are less sensitive to multicollinearity.
- For classical models (ARIMA, XGBoost without regularization, LightGBM default settings), you must clean collinear features first!



Recommended Quick Start

- For small projects: use pingouin or statsmodels VIF.
- For large pipelines: use **feature-engine** inside a preprocessing pipeline.